1. UDP packets are encapsulated with an IP header. An incoming segment has the source IP and hence, process can distinguish.

2. Yes.

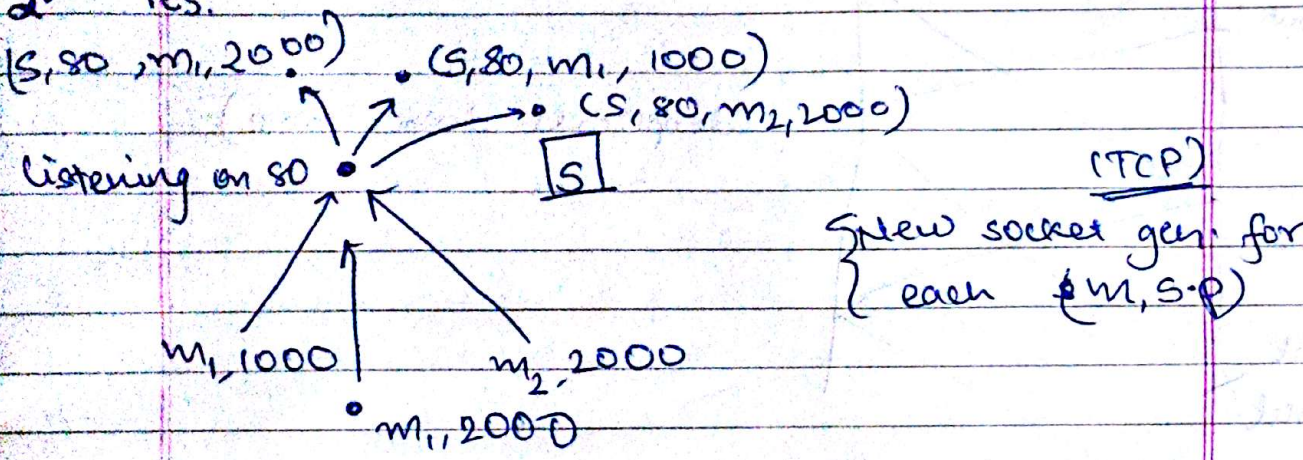

3. NO. Every new TCP connection from a (Source IP, source port) tuple creates & redirects to a new socket.

2 processes having same socket ⇒ same process. But since TCP is connection-oriented ...

4. (0 | 1 | 1 | 0)

3 2 1 0

It can capture error better than single-bit parity because single-bit parity is unaffected by a change each in o-e pairs. This has a better chance.

It adds an extra bit and can increased overhead, more than single-bit parity.

5 Bad forward / reverse channel, probabilistic loss

6.

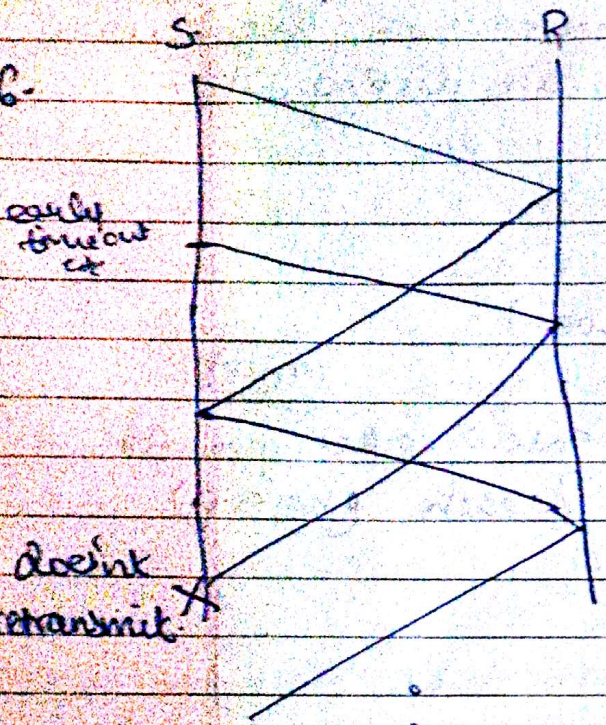

7.

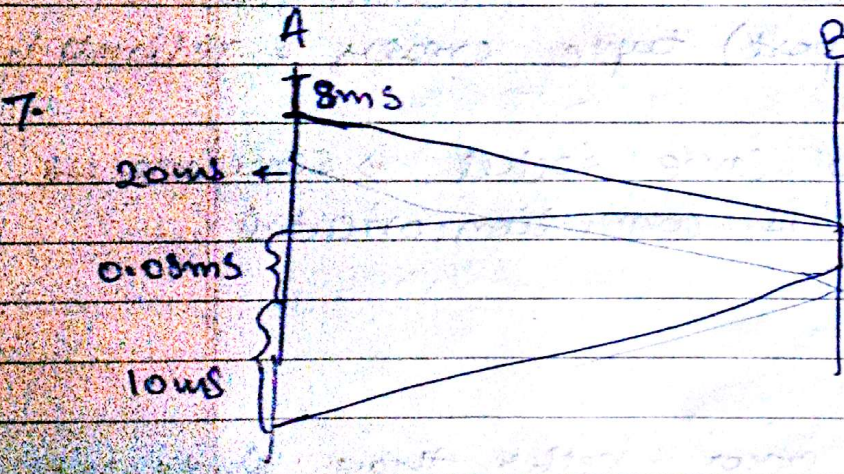

t.p. = data / time

= 1000 * 8 kbps / 38.08

= 210.08 kbps

8. Consider a pipeline (sliding window type) ck.

Best t.p. ⇒ ignore link delays? ∞ data?

≈ (8000 / 8.08) Kbps ≈ 1 Mbps

9. Maximum safe UDP payload is 508 bytes (max pkt 576 bytes − UDP header (8), IP (60))